

```

import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.model_selection import GridSearchCV

# Function to display comprehensive univariate analysis for
# categorical variables
def categorical_univariate_analysis(feature_name, data_series):
    # Frequency Distribution
    frequency_distribution = data_series.value_counts()

    # Display results
    print(f"\n----- Univariate Analysis for {feature_name} -----")
    print(f"Frequency Distribution:\n{frequency_distribution}\n")

    # Visualization
    plt.figure(figsize=(8, 5))
    sns.countplot(x=data_series,
order=data_series.value_counts().index) # Added order parameter
    plt.title(f'{feature_name} Distribution')
    plt.show()

def outlier_removal(feature_name, data_series):
    # IQR Method for Outlier Detection
    Q1 = data_series.quantile(0.25)
    Q3 = data_series.quantile(0.75)
    IQR = Q3 - Q1

    # Identify and remove outliers for the current feature
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    # Create a mask for non-outliers for the current feature
    outlier_mask = (df[feature_name] >= lower_bound) &
(df[feature_name] <= upper_bound)

    # Apply the mask to the DataFrame
    global df_no_outliers # Use global to modify the variable outside
the function
    if df_no_outliers is None:
        df_no_outliers = df[outlier_mask]
    else:
        df_no_outliers = df_no_outliers[outlier_mask]

```

```

# Display results
print(f"\n----- Univariate Analysis for {feature_name} -----")
print(f"IQR Method for Outlier Detection:")
print(f"Lower Bound: {lower_bound}")
print(f"Upper Bound: {upper_bound}")
print(f"Number of Outliers Removed: {len(df) -
len(df_no_outliers)}\n")

# Visualizations
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
sns.histplot(df_no_outliers[feature_name])
plt.title(f'{feature_name} Distribution (No Outliers)')

plt.subplot(1, 2, 2)
sns.boxplot(x=df_no_outliers[feature_name])
plt.title(f'{feature_name} Box Plot (No Outliers)')
plt.show()

# Assuming 'data.csv' is your dataset file
data_path = 'my_data.csv'

# Read the dataset
df = pd.read_csv(data_path)
df.info()

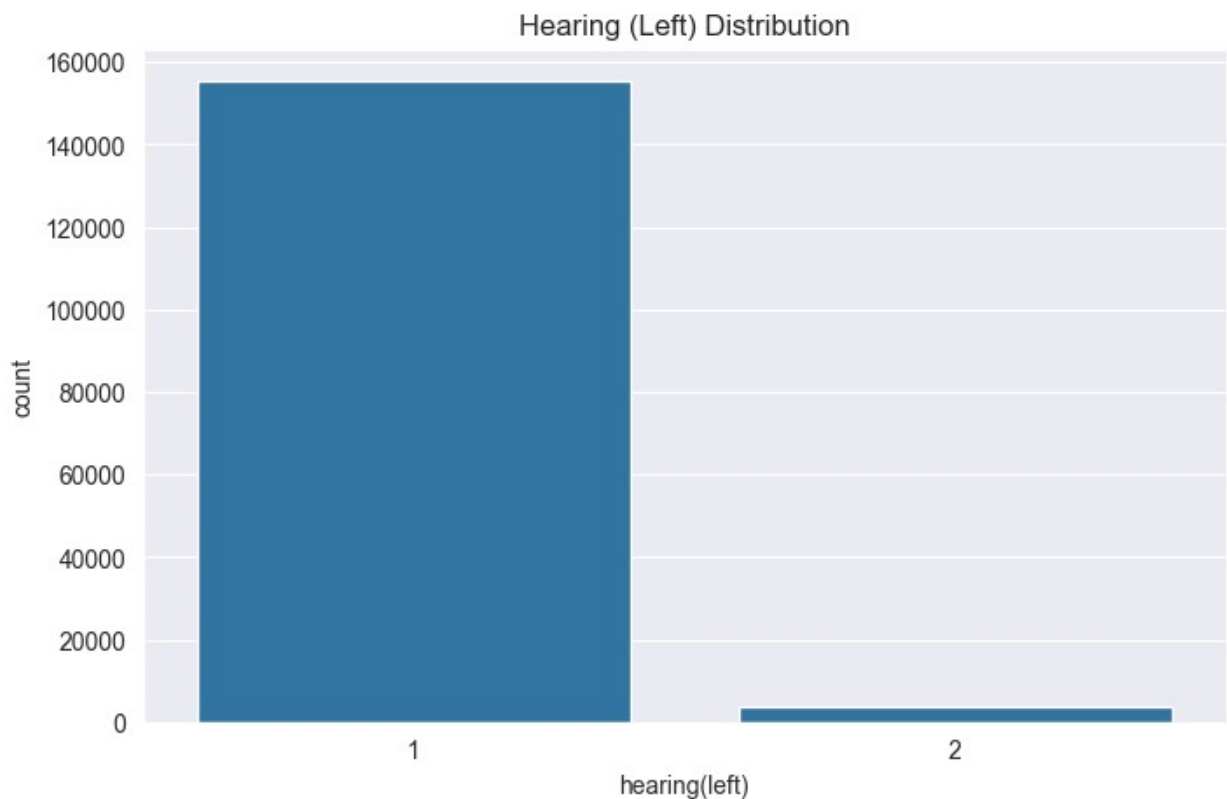
# DataFrame without outliers
df_no_outliers = None # Initialize to None

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 159256 entries, 0 to 159255
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     159256 non-null int64
1   ALT                    159256 non-null int64
2   AST                    159256 non-null int64
3   hearing(left)         159256 non-null int64
4   weight(kg)            159256 non-null int64
5   hearing(right)        159256 non-null int64
6   relaxation            159256 non-null int64
7   waist(cm)             159256 non-null float64
8   Cholesterol            159256 non-null int64
9   HDL                   159256 non-null int64
10  systolic               159256 non-null int64
11  smoking                159256 non-null int64
dtypes: float64(1), int64(11)
memory usage: 14.6 MB

```


```
# Apply the functions to each feature
categorical_univariate_analysis('Hearing (Left)', df['hearing(left)'])
```

```
----- Univariate Analysis for Hearing (Left) -----
Frequency Distribution:
hearing(left)
1      155438
2       3818
Name: count, dtype: int64
```



```
categorical_univariate_analysis('Hearing (Right)',
df['hearing(right)'])
```

```
----- Univariate Analysis for Hearing (Right) -----
Frequency Distribution:
hearing(right)
1      155526
2       3730
Name: count, dtype: int64
```



```
outlier_removal('ALT', df['ALT'])
```



```
----- Univariate Analysis for ALT -----
```

```
IQR Method for Outlier Detection:
```

```
Lower Bound: -8.0
```

```
Upper Bound: 56.0
```

```
Number of Outliers Removed: 6746
```



```
outlier_removal('AST', df['AST'])
```

```
C:\Users\DELL\AppData\Local\Temp\ipykernel_18984\232537705.py:19:
UserWarning: Boolean Series key will be reindexed to match DataFrame
index.
```

```
df_no_outliers = df_no_outliers[outlier_mask]
```



```
----- Univariate Analysis for AST -----
```

```
IQR Method for Outlier Detection:
```

```
Lower Bound: 6.5
```

```
Upper Bound: 42.5
```

```
Number of Outliers Removed: 10420
```



```
outlier_removal('weight(kg)', df['weight(kg)'])
```

```
C:\Users\DELL\AppData\Local\Temp\ipykernel_18984\232537705.py:19:
UserWarning: Boolean Series key will be reindexed to match DataFrame
index.
```

```
df_no_outliers = df_no_outliers[outlier_mask]
```



```
----- Univariate Analysis for weight(kg) -----
```

```
IQR Method for Outlier Detection:
```

```
Lower Bound: 37.5
```

```
Upper Bound: 97.5
```

```
Number of Outliers Removed: 12046
```



```
outlier_removal('relaxation', df['relaxation'])
```

```
C:\Users\DELL\AppData\Local\Temp\ipykernel_18984\232537705.py:19:
UserWarning: Boolean Series key will be reindexed to match DataFrame
index.
```

```
df_no_outliers = df_no_outliers[outlier_mask]
```



```
----- Univariate Analysis for relaxation -----
```

```
IQR Method for Outlier Detection:
```

```
Lower Bound: 52.0
```

```
Upper Bound: 100.0
```

```
Number of Outliers Removed: 12906
```



```
outlier_removal('waist(cm)', df['waist(cm)'])
```

```
C:\Users\DELL\AppData\Local\Temp\ipykernel_18984\232537705.py:19:
UserWarning: Boolean Series key will be reindexed to match DataFrame
index.
```

```
df_no_outliers = df_no_outliers[outlier_mask]
```



```
----- Univariate Analysis for waist(cm) -----
```

```
IQR Method for Outlier Detection:
```

```
Lower Bound: 59.0
```

```
Upper Bound: 107.0
```

```
Number of Outliers Removed: 13151
```

```
outlier_removal('Cholesterol', df['Cholesterol'])
```

```
C:\Users\DELL\AppData\Local\Temp\ipykernel_18984\232537705.py:19:  
UserWarning: Boolean Series key will be reindexed to match DataFrame  
index.
```

```
df_no_outliers = df_no_outliers[outlier_mask]
```



```
----- Univariate Analysis for Cholesterol -----
```

```
IQR Method for Outlier Detection:
```

```
Lower Bound: 112.0
```

```
Upper Bound: 280.0
```

```
Number of Outliers Removed: 13518
```



```
outlier_removal('HDL', df['HDL'])
```

```
C:\Users\DELL\AppData\Local\Temp\ipykernel_18984\232537705.py:19:
UserWarning: Boolean Series key will be reindexed to match DataFrame
index.
```

```
df_no_outliers = df_no_outliers[outlier_mask]
```



```
----- Univariate Analysis for HDL -----
```

```
IQR Method for Outlier Detection:
```

```
Lower Bound: 16.5
```

```
Upper Bound: 92.5
```

```
Number of Outliers Removed: 15362
```



```
outlier_removal('systolic', df['systolic'])
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_18984\232537705.py:19:
UserWarning: Boolean Series key will be reindexed to match DataFrame
index.

```
df_no_outliers = df_no_outliers[outlier_mask]
```



----- Univariate Analysis for systolic -----

IQR Method for Outlier Detection:

Lower Bound: 90.0

Upper Bound: 154.0

Number of Outliers Removed: 16543



```
categorical_univariate_analysis('smoking', df['smoking'])
```

```
----- Univariate Analysis for smoking -----
```

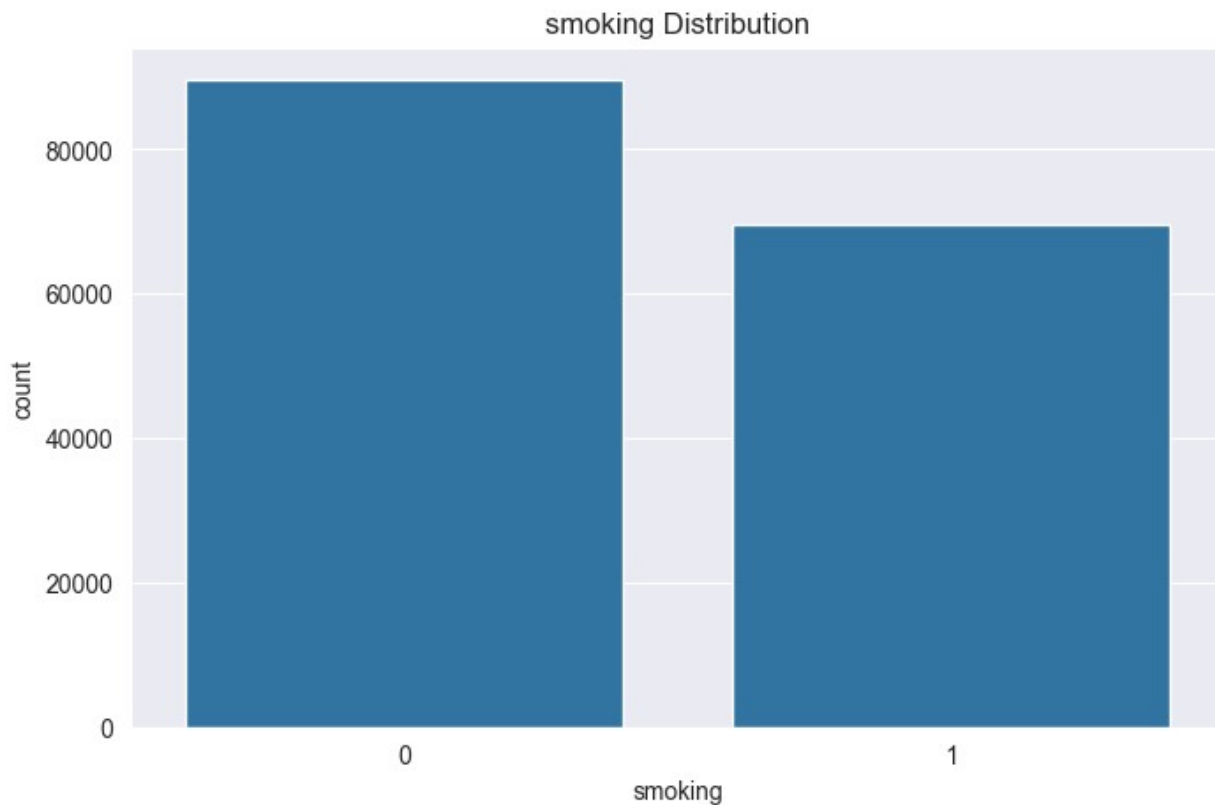
```
Frequency Distribution:
```

```
smoking
```

```
0      89603
```

```
1      69653
```

```
Name: count, dtype: int64
```



```
df_no_outliers.describe()
```

	id	ALT	AST	hearing(left)	\
count	142713.000000	142713.000000	142713.000000	142713.000000	
mean	79659.734208	23.992713	24.071290	1.024959	
std	45975.327633	10.632297	6.154968	0.156001	
min	0.000000	1.000000	7.000000	1.000000	
25%	39882.000000	16.000000	20.000000	1.000000	
50%	79641.000000	21.000000	23.000000	1.000000	
75%	119472.000000	30.000000	28.000000	1.000000	
max	159255.000000	56.000000	42.000000	2.000000	

	weight(kg)	hearing(right)	relaxation	waist(cm)	\
count	142713.000000	142713.000000	142713.000000	142713.000000	
mean	66.363898	1.024539	76.493059	82.493755	
std	11.663062	0.154715	8.651413	8.443666	
min	40.000000	1.000000	52.000000	59.000000	
25%	60.000000	1.000000	70.000000	77.000000	
50%	65.000000	1.000000	77.000000	83.000000	
75%	75.000000	1.000000	82.000000	88.000000	
max	95.000000	2.000000	100.000000	107.000000	

	Cholesterol	HDL	systolic	smoking
count	142713.000000	142713.000000	142713.000000	142713.000000
mean	195.355469	55.64711	121.938303	0.430991

std	28.124415	13.30671	12.149263	0.495217
min	112.000000	22.00000	90.000000	0.000000
25%	175.000000	45.00000	113.000000	0.000000
50%	196.000000	54.00000	120.000000	0.000000
75%	217.000000	64.00000	130.000000	1.000000
max	280.000000	92.00000	154.000000	1.000000

```
numeric_columns_to_normalize = ['ALT', 'AST', 'weight(kg)',
                                'waist(cm)', 'Cholesterol', 'HDL', 'relaxation', 'systolic']
```

```
# Create a copy of the DataFrame to avoid modifying the original data
df_normalized_zscore = df_no_outliers.copy()
```

```
# Apply Z-Score Normalization to selected columns
```

```
scaler_zscore = StandardScaler()
df_normalized_zscore[numeric_columns_to_normalize] =
scaler_zscore.fit_transform(
    df_no_outliers[numeric_columns_to_normalize])
df_normalized_zscore.reset_index(drop=True, inplace=True)
```

```
# Display the first few rows of the normalized DataFrame
```

```
print("Z-Score Normalized DataFrame:")
print(df_normalized_zscore)
```

Z-Score Normalized DataFrame:

	id	ALT	AST	hearing(left)	weight(kg)
hearing(right) \					
0	0	0.094739	-0.336524	1	-0.545647
1					
1	1	-0.093368	0.475830	2	-0.116942
2					
2	2	0.659059	0.475830	1	0.740469
1					
3	3	0.282846	-0.661466	1	2.455290
1					
4	4	-1.033902	-0.823937	1	-0.545647
1					
...
...					
142708	159251	0.188792	0.150888	1	-1.831764
1					
142709	159252	-0.375528	-0.498995	1	0.740469
1					
142710	159253	-1.410115	-1.473821	1	-1.403058
1					
142711	159254	-0.657688	-0.336524	1	0.740469
1					
142712	159255	-0.751742	-0.498995	1	-1.831764
1					

	relaxation	waist(cm)	Cholesterol	HDL	systolic
smoking					
0	1.214481	-0.176909	-0.830437	-1.175885	1.075106
1					
1	0.752127	0.770550	-0.048196	0.101670	1.980514
0					
2	-0.172580	-0.176909	-0.617098	-0.800134	-0.324161
1					
3	1.330070	2.665468	-0.545985	-1.326186	0.745867
0					
4	-0.056992	-0.236125	-1.434896	-0.875284	-0.077232
1					
...
.					
142708	0.405362	-1.598098	1.516287	1.228925	0.416627
0					
142709	0.405362	-0.058477	0.627376	0.627722	-0.159541
0					
142710	-0.750523	-1.953395	-0.225978	2.356180	-0.653400
0					
142711	1.561246	1.125847	-1.079332	-0.048631	-0.077232
1					
142712	1.214481	-0.721698	-0.332647	2.356180	0.252008
0					

[142713 rows x 12 columns]

```
# Dropping hearing and systolic and waist due to high co relation with
hearing and relaxation and weight
df_drop = df_normalized_zscore.drop(columns=['hearing(left)',
'systolic', 'waist(cm)', 'ALT'])
```

We either drop these features due to strong Co-Relation with their other features or use the dataframe from PCA both produced dimensionality reduction

```
# Assuming 'df_normalized_zscore' is your DataFrame with outliers
removed and normalized
numeric_columns_for_pca = ['ALT', 'AST', 'weight(kg)', 'relaxation',
'waist(cm)', 'Cholesterol', 'HDL', 'systolic']

# Standardize the data (important for PCA)
features_standardized_for_pca =
scaler_zscore.fit_transform(df_normalized_zscore[numeric_columns_for_p
ca])

# Apply PCA for dimensionality reduction
pca_for_replacement = PCA()
principal_components_for_replacement =
pca_for_replacement.fit_transform(features_standardized_for_pca)
```

```
# Variance explained by each principal component
explained_variance_ratio =
pca_for_replacement.explained_variance_ratio_

# The variable 'principal_components_for_replacement' contains the
transformed data
print("Principal Components:")
print(pd.DataFrame(principal_components_for_replacement,
                    columns=[f'PC{i + 1}' for i in
range(len(numeric_columns_for_pca))]))
```

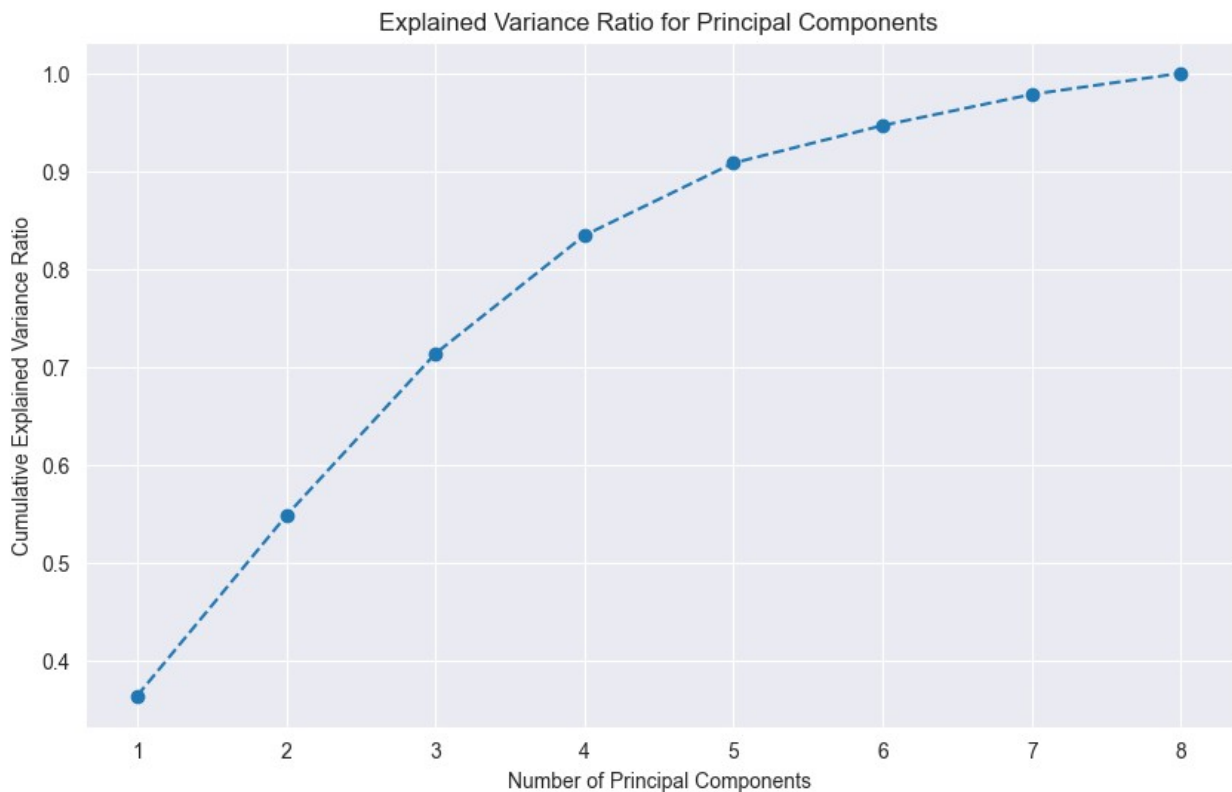
Principal Components:

	PC1	PC2	PC3	PC4	PC5	PC6
0	0.681363	1.156628	-0.712029	1.305247	-0.986761	-0.277206
1	1.263529	1.473325	0.292802	0.499670	0.211427	0.764212
2	0.697273	-1.001847	-0.274295	0.625430	-0.200735	-0.260789
3	3.462270	0.289643	-2.209468	-0.574935	0.709304	0.002516
4	-0.849491	0.219042	-1.696362	1.002570	-0.408563	0.246772
...
142708	-1.548713	1.315786	2.246320	-0.237603	-0.745740	-0.489491
142709	-0.034539	0.527167	-0.058236	-0.942857	0.515105	-0.272927
142710	-3.765110	0.961048	0.218436	-0.184108	0.988913	-0.551913
142711	0.984868	0.800524	-1.291138	0.399962	0.869098	0.102761
142712	-1.921453	2.090212	1.020938	0.597234	1.062709	-0.366772
PC8						
0	-0.293692					
1	-0.280173					
2	0.550928					
3	-0.275432					
4	-0.225743					
...	...					
142708	-0.129784					
142709	0.450742					
142710	0.213407					
142711	-0.504714					
142712	-0.992068					


```
[142713 rows x 8 columns]
```

```
# Plotting the explained variance ratio
```


```
plt.figure(figsize=(10, 6))
plt.plot(range(1, len(explained_variance_ratio) + 1),
         explained_variance_ratio.cumsum(), marker='o', linestyle='--')
plt.title('Explained Variance Ratio for Principal Components')
plt.xlabel('Number of Principal Components')
plt.ylabel('Cumulative Explained Variance Ratio')
plt.grid(True)
plt.show()
```



We deduced from the variance ratio that after the pc6 no huge variance occurs

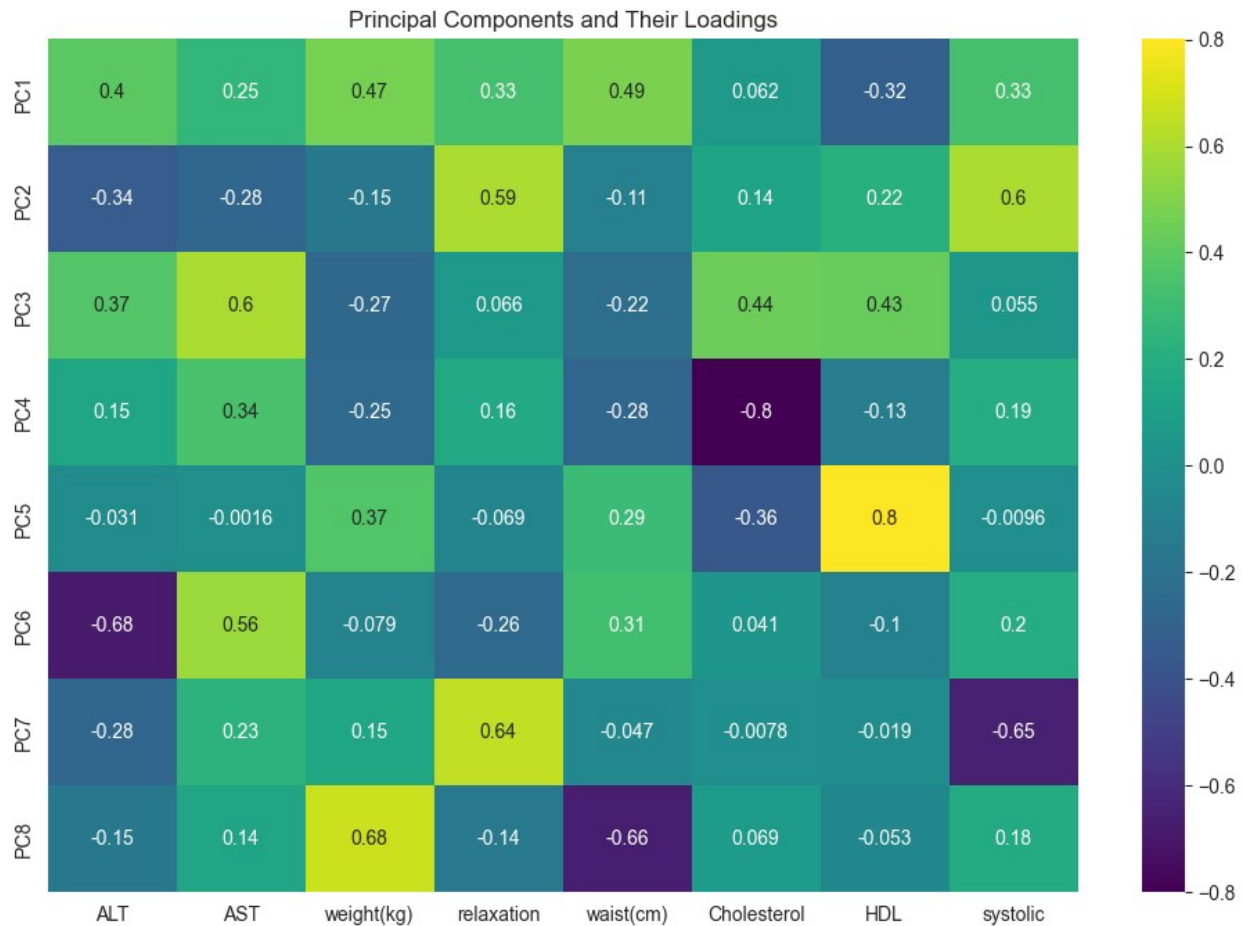
```
# Scree Plot
```

```
eigenvalues = pca_for_replacement.singular_values_ ** 2
plt.plot(range(1, len(eigenvalues) + 1), eigenvalues, marker='o')
plt.title('Scree Plot')
plt.xlabel('Principal Component')
plt.ylabel('Eigenvalue')
plt.show()
```



Here the scree plot confirmed that we should take the first 6 principal components

```
# Heatmap for loadings
plt.figure(figsize=(12, 8))
sns.heatmap(pca_for_replacement.components_, cmap='viridis',
            annot=True, xticklabels=numeric_columns_for_pca,
            yticklabels=[f'PC{i + 1}' for i in
range(len(numeric_columns_for_pca))])
plt.title('Principal Components and Their Loadings')
plt.show()
```



Now we will produce the final dataframe to start the ensemble methods with 6 principal components

```
# Assuming 'df_normalized_zscore' is your DataFrame with outliers
removed and normalized
numeric_columns_for_pca = ['ALT', 'AST', 'weight(kg)', 'relaxation',
'waist(cm)', 'Cholesterol', 'HDL', 'systolic']

# Standardize the data (important for PCA)
scaler_for_pca = StandardScaler()
features_standardized_for_pca =
scaler_for_pca.fit_transform(df_normalized_zscore[numeric_columns_for_
pca])

# Apply PCA for dimensionality reduction
pca_for_replacement = PCA(n_components=6) # Set the number of
components to 6
principal_components_for_replacement =
pca_for_replacement.fit_transform(features_standardized_for_pca)

# Create a DataFrame with the first two columns of 'hearing' from
df_normalized_zscore
```

```
final_data = df_normalized_zscore[['hearing(left)',
'hearing(right)']].copy()

# Add the 6 principal components to final_data
final_data = pd.concat(
    [final_data, pd.DataFrame(principal_components_for_replacement,
columns=[f'PC{i + 1}' for i in range(6)])], axis=1)

# Add the 'smoking' column from df_normalized_zscore
final_data['smoking'] = df_normalized_zscore['smoking']
# Display the final DataFrame
print("Final DataFrame:")
print(final_data)
```

```

142711  0.869098  0.102761      1
142712  1.062709 -0.366772      0

[142713 rows x 9 columns]

# Assuming 'final_data' is your DataFrame
X = final_data.drop(columns=['hearing(left)', 'hearing(right)',
                             'smoking'])
y = final_data['smoking']

# Split the data into training (70%), validation (15%), and testing
(15%) sets
X_train, X_temp, y_train, y_temp = train_test_split(X, y,
                                                    test_size=0.3, random_state=42)
X_valid, X_test, y_valid, y_test = train_test_split(X_temp, y_temp,
                                                    test_size=0.5, random_state=42)
y_train = np.ravel(y_train)
y_val = np.ravel(y_valid)
y_test = np.ravel(y_test)

```

Now we splitted the data to 70% training 15% validation 15% testing

```

class BaggingClassifier:
    def __init__(self, n_estimators=50, max_features=0.7,
                 max_depth=None):
        self.n_estimators = n_estimators
        self.max_features = max_features
        self.max_depth = max_depth
        self.estimators = []

    def fit(self, X, y):
        for _ in range(self.n_estimators):
            indices = np.random.choice(len(X), len(X), replace=True)
            X_bootstrap = X.iloc[indices]
            y_bootstrap = y[indices]
            estimator =
DecisionTreeClassifier(max_features=self.max_features,
max_depth=self.max_depth).fit(
            X_bootstrap, y_bootstrap)
            self.estimators.append(estimator)

    def predict(self, X):
        # Make predictions using all the base classifiers
        predictions = [estimator.predict(X) for estimator in
self.estimators]
        # Aggregate predictions using majority voting
        majority_votes = np.apply_along_axis(lambda x:
np.bincount(x).argmax(), axis=0, arr=predictions)

        return majority_votes

```

```

def get_params(self, deep=True):
    return {
        'n_estimators': self.n_estimators,
        'max_features': self.max_features,
        'max_depth': self.max_depth
    }

def set_params(self, **params):
    for param, value in params.items():
        setattr(self, param, value)
    return self

bagging_classifier_basic = BaggingClassifier()
# Fit BaggingClassifier on training data
bagging_classifier_basic.fit(X_train, y_train)

# Predict on validation set
y_pred = bagging_classifier_basic.predict(X_valid)
# Evaluate accuracy
accuracy = accuracy_score(y_valid, y_pred)
print(f"Accuracy on the validation set (Bagging): {accuracy:.2%}")

Accuracy on the validation set (Bagging): 64.87%

class AdaBoostClassifier:
    def __init__(self, n_estimators=50, learning_rate=1.0):
        self.n_estimators = n_estimators
        self.learning_rate = learning_rate
        self.estimators = []
        self.weights = []

    def fit(self, X, y):
        # Initialize Equal Weights
        self.weights = np.ones(len(X)) / len(X)
        for _ in range(self.n_estimators):
            # Train a decision stump
            estimator = DecisionTreeClassifier(max_depth=1).fit(X, y,
sample_weight=self.weights)

            # Compute Error rate
            predictions = estimator.predict(X)
            incorrect = (predictions != y)
            error_rate = np.dot(self.weights, incorrect)

            # Compute Alpha_t
            alpha = self.learning_rate * np.log((1 - error_rate) /
error_rate)
            self.estimators.append((estimator, alpha))

            # Reweighting

```

```

        self.weights *= np.exp(-y * alpha * predictions)

        # Normalization >> SUMMATION = 1
        self.weights /= np.sum(self.weights)

    def predict(self, X):
        predictions = np.empty(len(X))
        for estimator, alpha in self.estimators:
            predictions += alpha * estimator.predict(X)
        return np.sign(predictions)

    def get_params(self, deep=True):
        return {'n_estimators': self.n_estimators, 'learning_rate':
self.learning_rate}

    def set_params(self, **params):
        if 'n_estimators' in params:
            self.n_estimators = params['n_estimators']
        if 'learning_rate' in params:
            self.learning_rate = params['learning_rate']
        return self

boosting_classifier_basic = AdaBoostClassifier()

# Fit BaggingClassifier on training data
boosting_classifier_basic.fit(X_train, y_train)

# Predict on validation set
y_pred = boosting_classifier_basic.predict(X_valid)
# Evaluate accuracy
accuracy = accuracy_score(y_valid, y_pred)
print(f"Accuracy on the validation set (Boosting): {accuracy:.2%}")

Accuracy on the validation set (Boosting): 53.17%

class RandomForestClassifier:
    def __init__(self, n_estimators=500, max_features='sqrt',
max_depth=20, min_samples_split=10, min_samples_leaf=1):
        self.n_estimators = n_estimators
        self.max_features = max_features
        self.min_samples_split = min_samples_split
        self.min_samples_leaf = min_samples_leaf
        self.max_depth = max_depth
        self.estimators = []

    def fit(self, X, y):
        for _ in range(self.n_estimators):
            indices = np.random.choice(len(X), len(X), replace=True)
            X_bootstrap = X.iloc[indices]
            y_bootstrap = y[indices]
            estimator =

```

```

DecisionTreeClassifier(max_features=self.max_features,
max_depth=self.max_depth, min_samples_split = self.min_samples_split,
min_samples_leaf = self.min_samples_leaf).fit(X_bootstrap,
y_bootstrap)
        self.estimators.append(estimator)

    def predict(self, X):
        # Make predictions using all the base classifiers
        predictions = [estimator.predict(X) for estimator in
self.estimators]
        # Aggregate predictions using majority voting
        majority_votes = np.apply_along_axis(lambda x:
np.bincount(x).argmax(), axis=0, arr=predictions)
        return majority_votes

    def get_params(self, deep=True):
        return {
            'n_estimators': self.n_estimators,
            'max_features': self.max_features,
            'min_samples_split': self.min_samples_split,
            'min_samples_leaf': self.min_samples_leaf,
            'max_depth': self.max_depth
        }

    def set_params(self, **params):
        for param, value in params.items():
            setattr(self, param, value)
        return self

random_forest_classifier_basic = RandomForestClassifier()

# Fit BaggingClassifier on training data
random_forest_classifier_basic.fit(X_train, y_train)

# Predict on validation set
y_pred = random_forest_classifier_basic.predict(X_valid)
# Evaluate accuracy
accuracy = accuracy_score(y_valid, y_pred)
print(f"Accuracy on the validation set (Random Forest):
{accuracy:.2%}")

Accuracy on the validation set (Random Forest): 66.89%

from sklearn.metrics import make_scorer
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV

scoring = make_scorer(accuracy_score) # Use accuracy as the scoring
metric
bagging_classifier = BaggingClassifier()
param_distributions = {

```



```

        'n_estimators': [50, 100, 200],
        'max_features': [0.5, 0.7, 0.9],
        'max_depth': [None, 5, 10]
    }
    randomized_search = RandomizedSearchCV(bagging_classifier,
    param_distributions, n_iter=1, cv=5, n_jobs=-1,
                                     random_state=42,
    scoring=scoring)
    randomized_search.fit(X_train, y_train)
    best_bagging_model = randomized_search.best_estimator_

    y_pred = best_bagging_model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)

    print("\nBest hyperparameters for Bagging Classifier (Randomized
    Search):")
    print(best_bagging_model.get_params())
    print(f"Bagging Classifier Randomized Search Accuracy:
    {accuracy:.2%}")

```

Best hyperparameters for Bagging Classifier (Randomized Search):
 {'n_estimators': 200, 'max_features': 0.9, 'max_depth': None}
 Bagging Classifier Randomized Search Accuracy: 65.71%

```

    param_distributions = {
        'n_estimators': [20, 100, 200, 300],
        'learning_rate': [0.5, 1.0]
    }
    adaBoost_classifier = AdaBoostClassifier()
    randomized_search = RandomizedSearchCV(adaBoost_classifier,
    param_distributions, n_iter=1, cv=5, n_jobs=-1,
                                     random_state=42,
    scoring=scoring)
    randomized_search.fit(X_train, y_train)
    best_boosting_model = randomized_search.best_estimator_

    y_pred = best_boosting_model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)

    print("\nBest hyperparameters for AdaBoost Classifier (Randomized
    Search):")
    print(best_boosting_model.get_params())
    print(f"Boosting Classifier Randomized Search Accuracy:
    {accuracy:.2%}")

```

Best hyperparameters for AdaBoost Classifier (Randomized Search):
 {'n_estimators': 100, 'learning_rate': 0.5}
 Boosting Classifier Randomized Search Accuracy: 43.62%

```

param_grid = {
    'n_estimators': [150, 250, 350, 450],
    'max_depth': [15, 25, 35, 45],
    'min_samples_split': [2, 4, 8],
    'min_samples_leaf': [1, 2, 4]
}
randomForest_classifier = RandomForestClassifier()
randomized_search = RandomizedSearchCV(randomForest_classifier,
param_distributions, n_iter=1, cv=5, n_jobs=-1,
                                     random_state=42,
scoring=scoring)
randomized_search.fit(X_train, y_train)
best_randomForest_model = randomized_search.best_estimator_

y_pred = best_randomForest_model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)

print("\nBest hyperparameters for Random Forest Classifier (Randomized
Search):")
print(best_randomForest_model.get_params())
print(f"Random Forest Classifier Randomized Search Accuracy:
{accuracy:.2%}")

Best hyperparameters for Random Forest Classifier (Randomized Search):
{'n_estimators': 100, 'max_features': 'sqrt', 'min_samples_split': 10,
'min_samples_leaf': 1, 'max_depth': 20}
Random Forest Classifier Randomized Search Accuracy: 66.82%

```